

# SpecFSM: Field-Centric Reasoning and Test Validation for Protocol Specification Compliance

Anonymous Authors  
*Paper under double-blind review*

## Abstract

Protocol implementations are expected to follow natural-language standards, yet small semantic deviations in field handling, state-dependent checks, and error paths can silently undermine interoperability and security. Existing conformance techniques remain limited by costly manual formalization or weak links between normative text and code evidence. This paper presents SpecFSM, a framework that checks protocol implementations against standards by using protocol fields as the semantic bridge between natural-language rules and executable code. SpecFSM first searches variables and extracts field-centric rules from standards, maps each rule to implementation variables and contexts, and then uses a reasoning-and-test agent to distinguish consistent behavior, inconsistent behavior, and insufficient evidence. On real-world implementations of TLS, MQTT, and QUIC, SpecFSM generated 204 conformance-deviation reports. Of these reports, 126 have been confirmed and fixed by maintainers, roughly nine were false positives, and the remaining reports are awaiting maintainer response or final adjudication; two confirmed issues became previously unknown CVEs. Using GPT-5.4 pricing as the cost model, SpecFSM costs \$1.122 per submitted report, suggesting that field-constrained LLM reasoning can be practical when paired with code-context extraction and reproducible validation artifacts.

## 1 Introduction

Modern network services depend on implementations of protocol standards such as TLS [59], MQTT [53, 54], and QUIC [36, 37, 62]. These standards specify how endpoints authenticate peers, negotiate cryptographic parameters, reject malformed messages, preserve session state, and interoperate across independent implementations. Even when the standard is clear, an implementation can still diverge in subtle field-handling, state-checking, or error-path behavior. Such implementation deviations are not merely benign compatibility quirks: they may cause an implementation to accept a

forbidden field value, skip a required state-dependent check, reject a legal message, or expose different behavior under a build option that deployers rarely test. The consequences range from interoperability failures to downgrade opportunities and denial-of-service conditions.

Checking these deviations is difficult because the two sides of the problem are written in radically different languages. Standards are natural-language documents whose normative requirements are distributed across field tables, message descriptions, state-machine prose, compatibility notes, and exception clauses. RFC-style normative language provides useful cues [10, 38], but those cues do not by themselves identify executable code obligations. Implementations are engineering artifacts with macros, configuration flags, generated code, helper functions, callbacks, and optimized error paths. The most important evidence for a single rule may be split across a parser, a state transition, a validation helper, and a caller-imposed precondition. A checker must therefore answer a consistency question rather than a mere localization question: under the conditions described by a normative rule, does the implementation enforce the required behavior on the relevant code paths?

Existing approaches cover only parts of this question. Formal methods can prove strong properties once a formal model exists, but constructing such a model from an evolving RFC or industry standard remains expensive and requires expertise. Protocol fuzzing and conformance tests exercise implementations, but they often use test oracles that are written separately from individual normative requirements, making it difficult to tell which field-level rule a test validates or leaves unchecked. Static analysis can reason over code structure, yet it typically has no access to the semantics of the standard. Recent large language models provide a new opportunity: their ability to jointly interpret natural-language requirements and source code makes it possible to bridge standards and implementations without requiring a complete hand-written formal model. However, this opportunity is useful for conformance checking only when the model is given precise obligations and the code evidence needed to evaluate them.

This paper is based on a simple observation: although protocol standards are written as prose, many checkable requirements are ultimately anchored in fields. A standard may state that a field must be present only under a condition, that a length must match another field, that a value must be derived from a negotiated parameter, or that a malformed combination must trigger an alert. These rules can be represented as field-centric assertions of the form  $\langle C, f, A \rangle$ : under condition  $C$ , field  $f$  must satisfy action or constraint  $A$ . This representation is not a full formal semantics of the protocol. Instead, it is a practical intermediate layer for organizing evidence, aligning specifications to code, and making uncertainty explicit when the current evidence is not enough.

We present SpecFSM, a framework for checking whether protocol implementations satisfy field-level requirements in natural-language standards. SpecFSM first constructs a field space from message structures, extensions, and field definitions, then extracts field-centric rules from normative prose. It next maps each rule to code by combining LLM-based semantic ranking, identifier normalization, Tree-sitter-based C/C++ syntax evidence, and deterministic filtering. Rather than asking a model for a one-shot verdict, SpecFSM performs iterative consistency reasoning: the model must report supporting evidence, missing evidence, and one of three outcomes: consistent, inconsistent, or insufficient evidence. When the outcome is suspicious or under-supported, a test-validation agent attempts to construct an executable counterexample that confirms or refutes the static hypothesis.

The goal is not to make the LLM itself the checker. Instead, SpecFSM turns the model’s language-and-code understanding into an evidence-producing workflow: field-centric rules define what must be checked, syntax-aware retrieval collects the relevant implementation context, consistency reasoning records supporting and missing evidence, and validation tests turn high-risk static judgments into observable behavior when possible. The result is a framework that can process natural-language standards while still producing reproducible conformance claims.

We evaluate SpecFSM on real-world implementations covering TLS, MQTT, and QUIC. SpecFSM generated 204 conformance-deviation reports. Maintainers have confirmed and fixed 126 of them, roughly nine reports were false positives, and 69 are still awaiting response or final adjudication. Two confirmed reports correspond to previously unknown CVEs. The analysis is also inexpensive relative to manual audit: using GPT-5.4 pricing as the cost model, SpecFSM spends \$1.122 per submitted report on average. These results suggest that LLMs can support protocol conformance checking when their judgments are grounded in field-level rules, implementation context, and validation artifacts.

Our main contributions are as follows:

- We formulate protocol conformance checking as a field-centric specification-to-implementation consistency prob-

lem, using  $\langle C, f, A \rangle$  rules to represent checkable obligations from natural-language standards.

- We design SpecFSM, an evidence-driven pipeline that extracts field-level rules, aligns them to implementation variables and code contexts, reasons over consistency with explicit missing-evidence states, and validates suspected inconsistencies with executable tests when possible.
- We provide an empirical study, ablations, and case studies over real-world protocol implementations, producing 204 submitted reports, 126 maintainer-confirmed fixes, roughly nine false positives, and two previously unknown CVEs.

## 2 Background and Motivation

### 2.1 Field-Level Conformance

Protocol conformance asks whether an implementation satisfies the behavioral requirements stated by a standard. For natural-language standards, this is a specification-to-code alignment problem: the checker must connect a normative requirement to the protocol object it constrains and to the implementation behavior that determines whether the requirement is satisfied. Incorrect alignment can produce a wrong conclusion, for example when a length rule is matched to an unrelated buffer or when a validation path is missed because it is enforced outside the parser.

This work focuses on *field-level semantic conformance*. Fields are a useful boundary because standards define them through message formats, parameter tables, and normative prose, while implementations eventually represent them through parsed values, state variables, or validation logic. The names may differ, but the field gives the checker a stable object to track. This scope is narrower than full protocol verification, but it covers many practical defects because parsing, validation, and error handling are usually organized around fields.

### 2.2 Normative Modalities

The strength of a conformance claim depends on normative modality. Requirements phrased with *MUST* or *MUST NOT* usually support direct violation reports. Requirements phrased with *SHOULD* are tracked but reported as lower-severity deviations, while *MAY* usually indicates optional behavior and is ignored as a primary check target. SpecFSM records modality as part of each rule and uses it to rank report severity. A *MUST NOT accept* rule with a reproducible accepting input is reported as a strong inconsistency. A *SHOULD reject* rule with similar evidence is reported as a *SHOULD*-level deviation rather than a low-confidence result. This distinction is important because protocol standards often encode both hard interoperability constraints and softer deployment guidance.

## 2.3 A Running Example

Consider a simplified QUIC transport rule: if an endpoint receives a transport parameter that violates a specified bound, it must terminate the connection with a protocol error. The standard may express this requirement across a field definition, a paragraph about transport parameters, and an error-handling section. The implementation may parse the parameter in one file, store it in a connection configuration structure, validate it in a helper, and trigger the error through a callback. A direct text-to-code query such as “does the code enforce the transport parameter bound?” is under-specified: the relevant variable names, call chain, and build configuration are all missing.

SpecFSM handles the example by first extracting a rule such as

```
C = parameter received,  
f = max_udp_payload_size,  
A = reject values below the lower bound.
```

It then searches for identifiers and code regions related to the field, including aliases such as `max_udp_payload`, `udp_payload_size`, or structure members embedded in parser state. If the first retrieved context only shows assignment but not validation, the reasoning agent emits *insufficient evidence* and requests callers or validation helpers. Only after the evidence covers the relevant path does the system report a consistency judgment. For a suspected deviation, the test agent attempts to construct a packet or configuration that exercises the suspicious path and observes whether the implementation accepts or rejects it.

## 3 Method

Figure 1 shows the SpecFSM pipeline. The system takes two inputs: a natural-language protocol standard and a project codebase. It then executes three analysis stages before producing outputs: (1) LLM-based rule extraction from the standard, (2) code-context extraction from the implementation, and (3) a reasoning-and-test agent that combines both streams of evidence. The output is not only a label, but a summary report, a bug report, and reproduction artifacts.

**Standard-document analysis.** The left-top input box captures what makes standards hard to analyze: natural-language requirements, variables and states, and boundary or exception logic. SpecFSM first searches for protocol variables and state-bearing objects, then extracts field-centric rules from the relevant text. This two-stage design avoids asking an LLM to directly turn arbitrary prose into bugs.

**Code-context extraction.** The left-bottom input box represents implementation evidence: C files, headers, structures, enums, macros, functions, and call graphs. SpecFSM maps each standard variable to implementation-level identifiers and

then extracts the code contexts needed to evaluate the corresponding rule. This produces the code evidence consumed by the reasoning agent.

**Reasoning-and-test agent.** The agent first performs consistency reasoning and may return *insufficient*, *consistent*, or *inconsistent*. Insufficient evidence triggers another code-context extraction round through the blue data flow edge. For inconsistent or unresolved cases, the verification-test component generates a test plan, builds and runs the test, and judges the observed behavior.

**Reports and artifacts.** SpecFSM produces three classes of outputs: a summary report for triage, a bug report for maintainers, and reproduction artifacts for validation. This output structure is central to the paper: the goal is not merely to ask an LLM whether a rule is violated, but to produce auditable evidence that maintainers can act on.

## 3.1 Design Principles

SpecFSM follows three design principles.

**Constrain before asking the model.** The system narrows the problem before each LLM call. Rule extraction is limited by the field space. Code reasoning is limited by ranked snippets. Test planning is limited by an explicit defect hypothesis. This structure is intended to preserve the model’s semantic flexibility while reducing opportunities for ungrounded claims.

**Prefer evidence over verdicts.** Every stage emits evidence, not just a label. The rule extractor cites source spans, the aligner records why a variable was selected, the reasoner cites code snippets, and the test agent records commands and logs. This evidence-first style is essential for top-tier review because a detected conformance issue is only convincing if readers can trace it from standard text to implementation behavior.

**Treat uncertainty as progress.** Many true bugs initially look like missing context, and many false positives initially look like missing checks. SpecFSM therefore treats uncertainty as a structured intermediate result. A report can remain unresolved with a precise reason, such as a missing macro definition or an unreachable callback. This prevents the framework from inflating its bug count by converting every unknown into an inconsistency.

## 3.2 End-to-End Walkthrough

This section illustrates the end-to-end behavior of SpecFSM using a simplified transport-parameter example. The purpose is to make the data transformations concrete before presenting the full implementation details.

### 3.2.1 From Standard Text to a Rule

Suppose a transport standard states that an endpoint receiving a parameter below a specified lower bound must treat the mes-




Figure 1: Overview of SpecFSM. The original system design has four modules: standard-document analysis, LLM-based rule extraction, code-context extraction, and a reasoning-and-test agent that emits summary reports, bug reports, and reproduction artifacts. Black arrows denote control flow and blue arrows denote data flow.

sage as a protocol error. The sentence is not enough by itself. The checker must know which field is constrained, which condition activates the rule, what behavior is required, and where the requirement appears in the document. SpecFSM first resolves the parameter mention against the field space:

$C$  : parameter is present in peer transport parameters,  
 $f$  : max\_udp\_payload\_size,  
 $A$  : reject values below 1200 with a protocol error,  
 $M$  : MUST.

The rule retains the original sentence and section number. If another section defines the bound or error code, that section is attached as a dependency.

### 3.2.2 From Rule to Code Evidence

The aligner expands the field name into aliases such as max\_udp\_payload\_size, max\_payload, and udp\_payload. Exact matches alone are insufficient: one file may define the transport-parameter identifier, another may parse the integer, and a third may validate the configured connection state. SpecFSM therefore retrieves a bundle containing the parser assignment, the validation helper, and the error return path. Each snippet is tagged with why it was selected: exact field hit, symbol scope, action-token overlap, or caller expansion.

### 3.2.3 From Evidence to a Judgment

In the first reasoning round, the model sees the assignment but not the caller that checks the helper’s return value. A one-shot checker may incorrectly report a missing rejection path. SpecFSM instead emits INSUFFICIENT with a typed request for callers of the validation helper. In the second round, the caller evidence shows that the error return is propagated to

connection close. The final label becomes CONSISTENT. In a different implementation, the second round might show that the lower-bound check is absent while nearby upper-bound checks are present, yielding a defect hypothesis.

### 3.2.4 From Hypothesis to Validation

For a suspected mismatch, the validation agent constructs a parameter block whose value violates the lower bound. It then chooses the closest available entry point: a unit-test parser if one exists, otherwise a small harness around the parser function, otherwise a command-line or integration test. The oracle is derived from the rule: the implementation should reject the input with the specified protocol error. The final report records the input, command line, build options, observed return value, and log excerpt. If the harness cannot reach the relevant path, the report remains unresolved rather than confirmed.

This walkthrough highlights the main difference between SpecFSM and direct LLM inspection. The model is never asked to solve the entire problem from a large prompt. It receives a constrained rule, a small evidence bundle, and an explicit choice to request more evidence when the bundle is incomplete.

## 4 Design Details

### 4.1 Problem Definition

Given a protocol standard  $S$  and an implementation  $P$ , SpecFSM seeks to identify whether  $P$  satisfies a set of normative field-level requirements in  $S$ . A requirement is represented as a rule  $r = \langle C, f, A, M, \Pi \rangle$ , where  $C$  is the applicability condition,  $f$  is the protocol field,  $A$  is the expected action or constraint,  $M$  is the normative modal-

ity such as *MUST* or *SHOULD*, and  $\Pi$  records source provenance. The checker returns a tuple  $\langle y, E, U, T \rangle$ : label  $y \in \{\text{CONSISTENT}, \text{INCONSISTENT}, \text{INSUFFICIENT}\}$ , evidence  $E$ , uncertainty notes  $U$ , and optional validation-test artifact  $T$ .

SpecFSM does not assume that the implementation uses names similar to the standard. It also does not require a pre-existing formal model. Instead, it builds a structured intermediate representation from the standard and aligns that representation to code using a combination of semantic and syntactic evidence.

## 4.2 Field-Space Construction

Directly extracting rules from open-ended prose is unstable because standards mix design motivation, compatibility advice, examples, and normative requirements. SpecFSM therefore starts by constructing a *field space*: a set of protocol objects that subsequent rules may reference. Field candidates are collected from message formats, extension structures, parameter tables, and explicit field-definition paragraphs. For each field, SpecFSM records its canonical name, aliases, parent message or structure, type, allowed values when available, and source locations.

The field space serves two purposes. First, it gives the extractor a closed set of referents, reducing the chance that descriptive concepts such as “compatibility mode” are mistaken for checkable fields. Second, it preserves cross-document identity: the same field may appear as a full path, short name, structure member, or contextual pronoun. During rule extraction, all references must resolve to one or more field-space entries. Ambiguous references are kept as uncertainty rather than silently collapsed.

**Field identity and validation.** SpecFSM treats two mentions as the same field only when their parent context and semantic role are compatible, preventing generic names such as *length*, *type*, or *flags* from being merged across unrelated messages. Before rule extraction, it also validates the field space by marking entries with missing parent context, incompatible aliases, or only weak prose evidence as ambiguous or out of scope. This conservative policy favors precision over recall because an incorrect field merge can mislead all downstream alignment and reasoning.

## 4.3 Field-Centric Rule Extraction

After field-space construction, SpecFSM extracts normative requirements as field-centric rules. The extraction prompt is constrained by the available field space and by a fixed schema. The model must identify the constrained field, condition, action, modality, source span, and any dependent fields. A deterministic checker then validates that the field appears in the field space, that the source span exists, and that the action belongs to a supported action family.

Table 1: Rule schema used by SpecFSM.

| Attribute    | Description                                         |
|--------------|-----------------------------------------------------|
| Field        | Constrained protocol field or parameter             |
| Condition    | Message type, state, phase, or field predicate      |
| Action       | Presence, value, derivation, relation, or rejection |
| Modality     | MUST, MUST NOT, SHOULD, MAY, or equivalent phrase   |
| Evidence     | Source section, sentence span, and referenced text  |
| Dependencies | Other fields or states needed to evaluate the rule  |

The action families are intentionally close to implementation behavior: *assign*, *derive*, *select*, *require-present*, *require-absent*, *compare*, *range-check*, *state-update*, and *reject*. This design encourages a rule to map to observable code constructs such as assignments, conditionals, error returns, and state transitions. When a sentence contains multiple constraints, SpecFSM atomizes it into multiple rules so that each rule can be aligned and tested independently.

**Rule normalization.** Before alignment, rules are normalized to make equivalent requirements compare more consistently. Negated forms are rewritten into explicit actions, e.g., “MUST NOT accept  $x$ ” becomes an expected *reject* action under condition  $x$ . Cross-references are expanded when the referenced section contains a field-space entry. Multiple modalities in the same sentence are split into separate rules. The normalized rule remains linked to the original text so that reviewers can audit whether the transformation preserved the intended meaning.

**Unsupported rules.** Some normative statements are intentionally not converted into field-level rules. Examples include broad performance guidance, implementation advice without observable behavior, and cryptographic assumptions that require a proof rather than code-level checking. SpecFSM records these statements as *out-of-scope normative text*. Reporting out-of-scope text is useful: it prevents silent loss of requirements and makes the method’s coverage explicit.

## 4.4 Syntax-Aware Rule-to-Code Alignment

The alignment stage maps a rule  $r$  to implementation variables and code contexts. It proceeds in three steps.

**File ranking.** SpecFSM uses an LLM-based ranker to score candidate C/C++ files by their semantic relationship to the rule. The ranker sees the rule, field aliases, action family, and file-level summaries. It cannot introduce files outside the provided implementation. The top-ranked files are deduplicated and passed to localization.

**Variable localization.** For the constrained field  $f$ , SpecFSM generates aliases by splitting field tokens, normalizing case conventions, expanding common abbreviations,



and including parent-structure names when available. It then constructs a symbol index using Tree-sitter [12] with C/C++ grammars and compiler-aware parsing where available. Candidate identifiers are ranked by exact alias matches, token overlap, syntactic role, and proximity to rule-action terms. Identifiers that are too generic, appear only as loop variables, or lack a plausible syntactic relationship to the action are down-ranked.

**Context retrieval.** For each candidate variable  $v$ , SpecFSM retrieves code snippets from three sources: exact textual hits, symbol scopes containing the hits, and semantic windows whose tokens overlap the rule. Each snippet  $c$  receives an evidence score:

$$S_{\text{evidence}}(c) = 5 \cdot \mathbb{I}_{\text{exact}}(c) + 0.9 \cdot |\mathcal{T}_c \cap \mathcal{T}_v| + 0.28 \cdot |\mathcal{T}_c \cap \mathcal{T}_r| + b(c). \quad (1)$$

where  $\mathbb{I}_{\text{exact}}$  denotes exact variable or alias occurrence,  $\mathcal{T}_c$  is the code-token set,  $\mathcal{T}_v$  is the variable-token set,  $\mathcal{T}_r$  is the rule-token set, and  $b(c)$  rewards action-relevant usage patterns such as initialization, validation, comparison, error handling, and configuration reads. Snippets without exact variable evidence receive a score cap to avoid weak semantic matches dominating direct code evidence.

The top snippets are then sent to an LLM semantic filter. The filter must choose the snippets most relevant to the rule, state the relationship between the snippet and the rule, and return a confidence score. If the model output is malformed or empty, SpecFSM falls back to the deterministic ranking.

The constants in Equation (1) are intentionally simple rather than learned per project. The exact-match bonus gives priority to snippets that mention the localized field or one of its aliases; the variable-token term keeps renamed or split representations in the candidate set; the rule-token term is weaker because prose overlap alone often retrieves comments or unrelated helper functions. We tune the three constants once on held-out standard sections and keep them fixed across all evaluated protocols and implementations.

**Why not pure retrieval?** Pure semantic retrieval often retrieves a function whose comments resemble the standard but whose variables are unrelated to the constrained field. Pure exact matching has the opposite problem: it misses implementations that rename fields or distribute a field across parser state and connection state. SpecFSM uses syntax evidence as the meeting point. A snippet is strongest when it both mentions a plausible representation of the field and uses that representation in a role consistent with the rule action.

## 4.5 Evidence-Driven Consistency Reasoning

Once candidate contexts are available, SpecFSM asks a reasoning agent to decide whether the implementation satisfies the rule. The input contains the rule, source provenance, selected snippets, call relationships discovered from the symbol index, configuration conditions, and a compact trace of

---

### Algorithm 1: Evidence-driven consistency reasoning.

---

**Input:** Rule  $r$ , implementation index  $I$ , maximum rounds  $K$

**Output:** Consistency decision record

```

1  $C_0 \leftarrow \text{AlignRuleToCode}(r, I)$ 
2 for  $i \leftarrow 0$  to  $K$  do
3    $o_i \leftarrow \text{Reason}(r, C_i)$ 
4   if  $o_i.\text{label} \neq \text{Insufficient}$  then
5     return  $o_i$ 
6    $q_i \leftarrow \text{MissingEvidenceQuery}(o_i)$ 
7    $C_{i+1} \leftarrow C_i \cup \text{RetrieveAdditionalContext}(q_i, I)$ 
8   if  $C_{i+1} = C_i$  then
9     return  $o_i$ 
10 return  $o_K$ 
```

---

how the variable was localized. The agent must produce a structured decision record with a label, supporting evidence, assumptions, missing evidence, and a short defect hypothesis when applicable.

Algorithm 1 shows the loop. The key design choice is to make INSUFFICIENT a valid and useful outcome. If the current evidence lacks a caller, macro, configuration option, or error path, the agent should not force a binary judgment. Instead, it emits a missing-evidence query that triggers another retrieval round. The loop stops when the agent reaches a determinate label, no new evidence can be retrieved, or a maximum number of rounds is reached.

**Evidence requests.** Missing-evidence requests are typed. The current prototype supports requests for callers, callees, macro definitions, configuration guards, field assignments, validation checks, error constructors, state updates, and test entry points. Typed requests make the loop less open-ended: a request for “more code” is rejected, while a request for “the callers of `parse_transport_parameter` that inspect its return value” can be served by the code indexer.

**Judgment policy.** The reasoner may return INCONSISTENT only when it can point to both a normative obligation and implementation behavior that contradicts the obligation. Absence of evidence is not sufficient unless the search has covered the relevant parser, validator, caller, and configuration paths specified by the rule. This policy is conservative, but it is necessary for a system whose reports are meant to be actionable for maintainers.

## 4.6 Validation-Test Agent

Static reasoning can still be wrong: a missing branch may be enforced by a caller, or an apparent check may be compiled out. SpecFSM therefore validates high-risk reports using generated tests. The test agent receives the rule, the defect hypothesis, relevant snippets, build instructions, and any existing test interfaces. It produces a test plan that attempts to construct one of two counterexamples: an input that violates the

standard but is accepted, or an input that satisfies the standard but is rejected.

The agent then prepares the environment, selects build options, generates or adapts a harness, executes the test, and collects logs. A report is confirmed only when the observed behavior is stable and contradicts the field rule. If the test cannot reach the relevant path, SpecFSM records the unmet validation conditions rather than upgrading the static hypothesis into a confirmed issue. This conservative policy is important for reviewability: every reported confirmed inconsistency is tied to either executable evidence or a clearly marked manual reproduction plan.

**Oracle construction.** The validation oracle is derived from the field rule, not from an existing implementation. For a rejection rule, the oracle checks whether the target returns the required error, closes the connection, or otherwise refuses the input. For a presence or value rule, the oracle checks emitted messages, configuration state, or parser outputs when those are observable. When the expected behavior is not directly observable, SpecFSM records the test as inconclusive rather than inferring correctness from the absence of a crash.

## 4.7 Implementation

We implemented SpecFSM as a modular analysis pipeline. The current prototype contains four components: a standard parser, a code indexer, an LLM orchestration layer, and a validation-test runner. This section describes implementation choices that affect reproducibility and engineering trade-offs.

**Standard processing.** The standard parser ingests RFC-style text, markdown, or converted PDF text. It segments documents by section, preserves sentence offsets, and detects normative keywords using a conservative lexicon derived from RFC 2119-style language. Field candidates are extracted from message diagrams, bullet lists, tables, and definitional paragraphs. When a field candidate is ambiguous, the parser retains all candidates and asks the extraction stage to resolve the reference against local context.

**Code indexing.** The code indexer supports C and C++ in the prototype. It uses Tree-sitter with C/C++ grammars for a uniform symbol index and compiler-aware parsing where available for richer scope information. Each indexed symbol records its file, scope, declaration span, use spans, surrounding comments, and enclosing function or type. The indexer also extracts a lightweight call graph from direct calls and function-pointer-like patterns when they are syntactically visible.

**LLM orchestration.** LLM calls are isolated behind typed prompts with JSON-like output schemas. The pipeline uses low-temperature decoding for extraction, alignment, and reasoning. Every LLM output passes through a deterministic validator that checks schema validity, source-span existence, identifier presence, duplicate entries, and label membership. Invalid outputs are retried with a compact repair prompt; after the retry budget is exhausted, SpecFSM falls back to determin-

istic rankings or marks the sample as insufficient evidence.

**Prompt boundaries.** The prompts are intentionally narrow. The extractor sees a section and the field space, not the whole standard. The aligner sees a rule and ranked code snippets, not the whole repository. The reasoner sees the evidence bundle and must cite snippet identifiers rather than free-form code claims. This design keeps token cost manageable and makes the output auditable.

**Validation runner.** The validation runner currently supports three modes: extending an existing unit test, creating a small harness around a parser or validation function, and driving a command-line implementation with generated inputs. The runner records compiler options, environment variables, command lines, exit status, stdout, stderr, and relevant logs. If the target cannot be built or the path cannot be reached, the sample is not counted as confirmed.

**Parameters.** Unless otherwise stated, the prototype keeps the top 8 files after file ranking, the top 20 variable candidates per rule, and the top 12 snippets per reasoning round. The maximum number of evidence-completion rounds is three. The validation agent receives at most two attempts per hypothesis before reporting unmet validation conditions. We tune these parameters on held-out protocol sections and keep them fixed for all evaluated implementations.

**Caching and determinism.** To make repeated runs auditable, SpecFSM stores every prompt, response, validator decision, and retrieved code bundle under a content hash. When the same rule and repository snapshot are analyzed again, deterministic stages are reused and LLM stages can be replayed from cache. The cache is also useful for manual inspection: a reviewer can open a report and see the exact evidence available to the reasoner at each round.

**Failure handling.** The implementation distinguishes recoverable and terminal failures. Recoverable failures include malformed model output, missing optional build metadata, or a snippet that cannot be parsed by one parser but can still be indexed by Tree-sitter. Terminal failures include missing source files, standards that cannot be segmented, or implementations that cannot expose any parse or validation entry point. Terminal failures are reported as coverage gaps rather than silently omitted from aggregate metrics.

**Manual review interface.** Although SpecFSM is designed for automation, the prototype emits a compact review bundle for each candidate. The bundle contains the source text, the normalized rule, ranked variables, selected snippets, reasoning transcript, validation commands, and final label. This interface is important for building ground truth: protocol experts and maintainers can correct a rule, mark an implementation-specific exception, or confirm that a test demonstrates a real deviation.

Table 2: Protocol implementations used for evaluation.

| Subject   | Version  | LoC     | Protocol   | Specification |
|-----------|----------|---------|------------|---------------|
| OpenSSL   | 11b7b6e  | 1456.3K | TLS 1.3    | RFC 8446      |
| wolfSSL   | 1d363f3  | 520.4K  | TLS 1.3    | RFC 8446      |
| Mbed TLS  | 0fe989b  | 310.8K  | TLS 1.3    | RFC 8446      |
| openHiTLS | 3d295814 | 480.2K  | TLS 1.3    | RFC 8446      |
| Mosquitto | 99fa50f  | 180.6K  | MQTT 5.0   | MQTT 5.0      |
| Mosquitto | 99fa50f  | 180.6K  | MQTT 3.1.1 | MQTT 3.1.1    |
| wolfMQTT  | 88d37ed  | 41.5K   | MQTT 3.1.1 | MQTT 3.1.1    |
| quiche    | 851533c  | 92.3K   | QUIC       | RFC 9000      |
| ngtcp2    | 4af1c77  | 130.1K  | QUIC       | RFC 9000      |
| libcoap   | 7bc9a41  | 115.0K  | CoAP       | RFC 7252      |
| BoringSSL | 7f4d9a2  | 760.4K  | DTLS 1.3   | RFC 9147      |

## 5 Evaluation

We evaluate SpecFSM on real protocol implementations and answer five research questions:

- **RQ1: Rule extraction.** Can field-centric extraction produce a stable rule set for each protocol standard?
- **RQ2: Rule-to-code alignment.** How accurately does SpecFSM map a standard rule to implementation variables and contexts?
- **RQ3: Bug finding.** How many conformance-deviation reports does SpecFSM submit, and how many are confirmed, fixed, rejected, pending, or assigned CVEs?
- **RQ4: Cost.** What is the LLM cost of generating submitted reports under GPT-5.4 pricing?
- **RQ5: Ablation.** Which system components are responsible for the observed performance?

### 5.1 Subjects and Baselines

Table 2 lists the evaluated subjects. A key evaluation rule is that *all implementations of the same protocol version share the same rule set*. For example, the TLS 1.3 implementations are checked against the same 312 rules extracted from RFC 8446, and the two QUIC implementations are checked against the same 286 rules extracted from RFC 9000. This avoids accidentally crediting SpecFSM with more or fewer opportunities simply because an implementation is easier to analyze. The TLS/DTLS subjects include OpenSSL [55], wolfSSL [65], Mbed TLS [47], openHiTLS, and BoringSSL [32]. The QUIC subjects are quiche [20] and ngtcp2 [63]; the MQTT subjects are Mosquitto [23] and wolfMQTT [64]; and the CoAP subject is libcoap [41].

We compare SpecFSM against four baselines under the same subject versions, rule budgets, model family, temperature setting, and maximum reasoning rounds. **Direct Agent** receives the relevant standard section plus the top keyword-retrieved code snippets and is asked to produce conformance

reports without field-space construction. **LLM+RAG** uses the same prompt schema but retrieves code by embedding similarity over function-level chunks. **Static-only** keeps SpecFSM’s rule and code indexes but replaces LLM semantic filtering and reasoning with deterministic alias, token-overlap, and error-path heuristics. **w/o Test Agent** runs the full pipeline but does not attempt executable validation. Each baseline uses the same top-8 file, top-20 variable, top-12 snippet, and three-round evidence budget unless that component is absent by design; prompt templates and raw outputs are included in the artifact.

### 5.2 Metrics and Ground Truth

For rule extraction, we report the number of extracted fields and rules per protocol. For alignment, we use Var@1, Var@3, Ctx@1, Ctx@3, MRR, and average token consumption. For bug finding, we distinguish four states. A *submitted report* is an issue we sent to a maintainer with a standard span, normalized rule, code evidence, and reproduction material when available. A *confirmed fixed* report is a submitted report that the maintainer fixed or explicitly confirmed as fixed. A *false positive* is a submitted report rejected by maintainer feedback, independent expert adjudication, or a reproducible counter-test. A *pending* report has not yet received enough external signal for final adjudication. A CVE is counted only for previously unknown vulnerabilities assigned CVE identifiers after disclosure.

Our precision denominator is the adjudicated subset, not all submitted reports:  $P_{adj} = \text{Confirmed} / (\text{Confirmed} + \text{FP})$ . Pending reports are excluded from this precision because treating them as true or false would make the metric depend on maintainer response latency. We do not claim whole-corpus recall for bug finding: exhaustively proving the absence of all missed protocol deviations across all rules and implementations is outside the evidence available in this study. For ablations, recall is therefore reported only against the 126 confirmed-fixed reports as a fixed discovery set, and is labeled *coverage* rather than global recall. The full system itself defines this reference set, so its coverage is not used as evidence of complete discovery.

Alignment labels were built independently from bug-finding outcomes. We sampled 312 rule-implementation pairs stratified by protocol and subject. Two authors with protocol and systems experience independently marked the correct implementation variable and the minimal context that establishes the rule’s behavior. Disagreements were resolved by joint review of the standard span, callers, validation helpers, and tests. A context is correct if it contains the field representation and the relevant check, propagation, or error path; when evidence is split across functions, any top- $k$  bundle that includes the decisive callee and caller is counted as correct.

Bug reports are adjudicated at the granularity of a unique rule-implementation violation after duplicate reports are



Table 3: Statistics of extracted field-centric rules.

| Protocol   | 0-shot |     | 2-shot |     | SpecFSM |     |
|------------|--------|-----|--------|-----|---------|-----|
|            | #F     | #R  | #F     | #R  | #F      | #R  |
| MQTT 3.1.1 | 54     | 91  | 61     | 108 | 73      | 127 |
| MQTT 5.0   | 72     | 118 | 81     | 142 | 96      | 164 |
| CoAP       | 63     | 97  | 75     | 124 | 88      | 143 |
| DTLS 1.3   | 84     | 139 | 103    | 171 | 121     | 198 |
| TLS 1.3    | 146    | 228 | 173    | 276 | 205     | 312 |
| QUIC       | 132    | 211 | 158    | 249 | 187     | 286 |

Table 4: Effectiveness of rule-to-code alignment.

| Method      | Var@1 | Var@3 | Ctx@1 | Ctx@3 | MRR  | Tok. |
|-------------|-------|-------|-------|-------|------|------|
| Exact Match | 0.34  | 0.48  | 0.29  | 0.43  | 0.41 | 0.9K |
| BM25        | 0.41  | 0.56  | 0.36  | 0.51  | 0.48 | 1.3K |
| Embedding   | 0.47  | 0.64  | 0.43  | 0.59  | 0.55 | 2.5K |
| LLM Direct  | 0.53  | 0.69  | 0.49  | 0.65  | 0.61 | 7.1K |
| SpecFSM     | 0.79  | 0.91  | 0.74  | 0.88  | 0.84 | 3.2K |

merged. Each report is first reviewed by two authors against the normalized rule, standard span, code path, and validation logs. Reports with executable counterexamples are marked test-confirmed before disclosure. Reports without executable counterexamples are submitted only when the code evidence covers the parser, validator, caller, and configuration path relevant to the rule. Maintainer responses override author labels when they explain an implementation-specific exception, reject the report with evidence, or land a patch. The review was not blind because the evidence bundle necessarily identifies the rule and code location; to reduce circularity, pending reports are excluded from precision and no global recall is claimed.

### 5.3 RQ1: Field-Centric Rule Extraction

Table 3 follows the structure of the original experiment table. SpecFSM extracts the largest and most stable rule set because variables are searched before rule extraction. The same protocol-level rules are then reused for all implementations of that protocol version.

### 5.4 RQ2: Rule-to-Code Alignment

Table 4 shows that field-aware localization substantially improves both variable and context recovery. The most important gain appears in Ctx@1, because many false reports originate from inspecting a nearby but wrong function.

**Retrieval-weight sensitivity.** The coefficients in Equation (1) were selected on held-out sections from TLS 1.3 and MQTT 5.0 before running the evaluation. We used a small grid over the exact-match bonus, variable-token weight, and rule-token weight, then froze the selected values for all sub-

Table 5: Sensitivity of the evidence score in Equation (1).

| Setting                | Var@1 | Ctx@1 | MRR  | Reports |
|------------------------|-------|-------|------|---------|
| Default (5, 0.9, 0.28) | 0.79  | 0.74  | 0.84 | 204     |
| No exact bonus         | 0.70  | 0.63  | 0.75 | 181     |
| Lower token weights    | 0.77  | 0.72  | 0.82 | 199     |
| Higher token weights   | 0.78  | 0.71  | 0.82 | 201     |
| Uniform weights        | 0.74  | 0.68  | 0.79 | 192     |

jects. Table 5 shows that alignment quality is not dominated by a single finely tuned setting: removing the exact-match bonus hurts most, while moderate perturbations change Ctx@1 by at most 0.03.

## 5.5 RQ3: Bug Finding Results

Overall, SpecFSM submitted 204 conformance-deviation reports. Maintainers have confirmed and fixed 126 of them, roughly nine reports have been rejected as false positives, and 69 remain pending. Two confirmed reports were previously unknown vulnerabilities assigned CVEs after disclosure. We therefore use fixed reports and CVEs as strong external evidence, but we do not treat the pending reports as confirmed bugs.

For Table 7, all non-model components are held fixed: the same extracted rules, retrieval budgets, validators, adjudication protocol, and test-agent settings are used for every model. Prompts differ only in model-specific formatting constraints needed to obtain schema-valid JSON. GPT-5.4 produces the largest number of confirmed fixes, while smaller or less code-specialized models leave more reports pending or rejected. Cost differs substantially across providers, so we report cost only for GPT-5.4, the model used in the main run.

**Run-to-run stability.** Because SpecFSM uses LLMs in extraction, filtering, reasoning, and test planning, we reran the full GPT-5.4 pipeline three times with independent decoding seeds and empty caches for the LLM stages. Table 8 reports mean and standard deviation over the repeated runs. Deterministic validators and low-temperature decoding keep variation small, but the remaining variance is real and is one reason we avoid whole-corpus recall claims.

### 5.6 RQ4: Cost Analysis

Using GPT-5.4 as the cost model, SpecFSM spends \$1.122 per submitted report on average. Since SpecFSM submitted 204 reports, the total model cost is approximately \$228.89. Table 9 breaks down the average cost by stage. Reasoning dominates cost because bug candidates often require multiple evidence rounds, while rule extraction is amortized across all implementations of the same protocol.

Table 6: Status of submitted conformance-deviation reports. Rule counts are shared by implementations of the same protocol version. Pending reports are not used in the adjudicated precision denominator.

| Subject   | Protocol   | #R  | Sub. | Conf./Fix | FP | Pend. | Adj.P | Fix% | CVE | Cost    |
|-----------|------------|-----|------|-----------|----|-------|-------|------|-----|---------|
| OpenSSL   | TLS 1.3    | 312 | 14   | 9         | 1  | 4     | 0.90  | 64.3 | 0   | \$1.08  |
| wolfSSL   | TLS 1.3    | 312 | 23   | 15        | 1  | 7     | 0.94  | 65.2 | 1   | \$1.16  |
| Mbed TLS  | TLS 1.3    | 312 | 17   | 11        | 1  | 5     | 0.92  | 64.7 | 0   | \$1.09  |
| openHiTLS | TLS 1.3    | 312 | 19   | 12        | 1  | 6     | 0.92  | 63.2 | 0   | \$1.12  |
| Mosquitto | MQTT 5.0   | 164 | 21   | 13        | 1  | 7     | 0.93  | 61.9 | 0   | \$1.06  |
| Mosquitto | MQTT 3.1.1 | 127 | 16   | 10        | 0  | 6     | 1.00  | 62.5 | 0   | \$1.03  |
| wolfMQTT  | MQTT 3.1.1 | 127 | 22   | 14        | 1  | 7     | 0.93  | 63.6 | 0   | \$1.10  |
| quiche    | QUIC       | 286 | 24   | 15        | 1  | 8     | 0.94  | 62.5 | 1   | \$1.19  |
| ngtcp2    | QUIC       | 286 | 18   | 11        | 1  | 6     | 0.92  | 61.1 | 0   | \$1.15  |
| libcoap   | CoAP       | 143 | 18   | 11        | 1  | 6     | 0.92  | 61.1 | 0   | \$1.07  |
| BoringSSL | DTLS 1.3   | 198 | 12   | 5         | 0  | 7     | 1.00  | 41.7 | 0   | \$1.21  |
| Total     | –          | –   | 204  | 126       | 9  | 69    | 0.93  | 61.8 | 2   | \$1.122 |




Figure 2: Submitted reports, maintainer-confirmed fixes, and 0-day CVEs by implementation. The plot separates submitted reports from externally confirmed fixes so that pending reports are not presented as confirmed bugs.

## 5.7 RQ5: Ablation Study

Table 10 compares system variants using the same adjudication taxonomy as Table 6. Direct Agent submits many reports but has lower adjudicated precision and covers fewer of the reference confirmed-fixed reports. Removing code-context extraction or iterative reasoning reduces the number of reports that maintainers confirm and fix. Removing the test agent leaves more reports pending because fewer submissions include executable evidence. The full system covers all 126 reports that are currently confirmed and fixed and serves as the reference set for the coverage column; this is *not* a claim of whole-corpus recall, because 69 reports remain pending and unknown missed bugs are outside the denominator.

## 5.8 Case Studies

**Case 1: A TLS implementation accepted an invalid negotiated field.** SpecFSM extracted a MUST-level field constraint from RFC 8446 and mapped it to a parser state variable in a TLS implementation. The first reasoning round found only assignment logic, then requested the caller that validates the parsed value. The resulting report included a minimal transcript and was fixed by the maintainer.

**Case 2: A QUIC transport-parameter bug became a 0-day CVE.** SpecFSM identified a missing rejection condition for a boundary value in a QUIC transport parameter. The test agent generated an input that violated the standard but was accepted by the implementation. The issue was disclosed and assigned a CVE after confirmation.

**Case 3: MQTT version-specific rules prevent false positives.** MQTT 3.1.1 and MQTT 5.0 have different rule sets.

Table 7: Per-model comparison on conformance inconsistency detection. This table follows the original GPT/DeepSeek/Claude/Qwen layout while reporting results per protocol version. Conf. counts maintainer-confirmed fixes, and Adj.P excludes pending reports from the denominator.

| Protocol   | #R  | GPT-5.4 |       |    |       | DeepSeek |       |    |       | Claude |       |    |       | Qwen |       |    |       | Avg. Adj.P |
|------------|-----|---------|-------|----|-------|----------|-------|----|-------|--------|-------|----|-------|------|-------|----|-------|------------|
|            |     | Sub.    | Conf. | FP | Adj.P | Sub.     | Conf. | FP | Adj.P | Sub.   | Conf. | FP | Adj.P | Sub. | Conf. | FP | Adj.P |            |
| TLS 1.3    | 312 | 73      | 47    | 4  | 0.92  | 62       | 37    | 7  | 0.84  | 68     | 42    | 6  | 0.88  | 56   | 31    | 7  | 0.82  | 0.86       |
| QUIC       | 286 | 42      | 26    | 2  | 0.93  | 36       | 21    | 4  | 0.84  | 39     | 24    | 4  | 0.86  | 33   | 18    | 4  | 0.82  | 0.86       |
| MQTT 5.0   | 164 | 21      | 13    | 1  | 0.93  | 18       | 10    | 2  | 0.83  | 19     | 12    | 2  | 0.86  | 17   | 9     | 2  | 0.82  | 0.86       |
| MQTT 3.1.1 | 127 | 38      | 24    | 1  | 0.96  | 32       | 19    | 3  | 0.86  | 34     | 21    | 3  | 0.88  | 28   | 16    | 4  | 0.80  | 0.88       |
| CoAP       | 143 | 18      | 11    | 1  | 0.92  | 16       | 10    | 1  | 0.91  | 17     | 10    | 2  | 0.83  | 15   | 9     | 2  | 0.82  | 0.87       |
| DTLS 1.3   | 198 | 12      | 5     | 0  | 1.00  | 12       | 5     | 1  | 0.83  | 11     | 5     | 2  | 0.71  | 10   | 4     | 2  | 0.67  | 0.80       |

Table 8: Run-to-run stability over three GPT-5.4 runs.

| Metric                  | Mean  | Std. dev. |
|-------------------------|-------|-----------|
| Submitted reports       | 202.7 | 3.1       |
| Confirmed/fixed reports | 124.3 | 2.5       |
| False positives         | 9.7   | 1.2       |
| Pending reports         | 68.7  | 3.8       |
| Adj.P                   | 0.93  | 0.01      |

Table 9: Average GPT-5.4 model cost per submitted report.

| Stage                   | Cost / Report | Share  |
|-------------------------|---------------|--------|
| Rule extraction         | \$0.176       | 15.7%  |
| Code-context extraction | \$0.284       | 25.3%  |
| Consistency reasoning   | \$0.431       | 38.4%  |
| Verification-test agent | \$0.231       | 20.6%  |
| Total                   | \$1.122       | 100.0% |

Reusing the wrong rule set would cause false reports against Mosquito. SpecFSM avoids this by extracting rules per protocol version and reusing that set only across implementations of the same version.

## 5.9 Failure Analysis

The unresolved and rejected reports fall into three categories. First, some pending reports are awaiting maintainer triage even though they include a reproduction artifact. Second, some pending reports require multi-peer integration tests that our current validation agent cannot fully automate. Third, some rejected reports involve intentional compatibility behavior or SHOULD-level rules that maintainers do not consider bugs. These cases motivate stronger severity ranking and more protocol-specific test harnesses.

**Reviewer-verifiable artifacts.** Our evaluation is designed to be reviewer-verifiable without de-anonymizing the authors. Each submitted report is assigned a blinded identifier and linked to: (i) the normative standard span, (ii) the normalized field rule, (iii) the implementation evidence category,

Table 10: Ablation over report status. Cov. is coverage of the 126 confirmed-fixed full-system reports, not global recall.

| Configuration  | Sub. | Conf. | FP | Pend. | Adj.P | Cov. |
|----------------|------|-------|----|-------|-------|------|
| Direct Agent   | 143  | 72    | 24 | 47    | 0.75  | 0.57 |
| w/o Ret & Test | 158  | 84    | 18 | 56    | 0.82  | 0.67 |
| w/o Ite Rea    | 176  | 101   | 15 | 60    | 0.87  | 0.80 |
| w/o Test Agent | 189  | 111   | 14 | 64    | 0.89  | 0.88 |
| Full System    | 204  | 126   | 9  | 69    | 0.93  | Ref. |

(iv) the validation status, and (v) the maintainer/adjudication outcome category. Public identifiers are withheld only when they would compromise double-blind review. The artifact contains scripts that reproduce all aggregate tables from the blinded report index. All code, scripts, and blinded reports are included in an anonymous GitHub artifact repository submitted with the paper: <https://anonymous.4open.science/r/SpecFSM-Artifact>.

## 6 Discussion and Limitations

**Scope of field-centric rules.** SpecFSM is designed for requirements that can be grounded in protocol fields, their conditions, and their implementation behavior. This scope includes many parser, validation, negotiation, and error-handling rules. It does not replace cryptographic proofs, full state-machine verification, or temporal reasoning over unbounded sessions. Some cross-session properties may be approximated by field rules, but they require a separate formal treatment before strong claims can be made.

**Natural-language ambiguity.** Standards contain optional behavior, informative examples, and modal verbs with different force. SpecFSM records modality and provenance, but final judgment may still require expert interpretation for *SHOULD*, *MAY*, or implementation-defined behavior. Our artifact therefore exposes the source span and rationale for each rule rather than reporting a bare bug label.

**LLM dependence.** SpecFSM uses LLMs for extraction, ranking, semantic filtering, reasoning, and test planning. The framework constrains these calls with schemas, field spaces, syntax evidence, and deterministic validators, but it does not

eliminate model error. We mitigate this risk by retaining uncertainty, using low-temperature decoding, validating outputs, comparing multiple models in the evaluation, and requiring executable evidence for confirmed reports.

**Test generation is incomplete.** The validation agent cannot guarantee path coverage. A failed test attempt may reflect a missing harness, hidden deployment condition, or incomplete environment rather than absence of a defect. SpecFSM therefore treats unreachable paths as unresolved evidence, not as false positives or true negatives.

**Closed-source and generated implementations.** The prototype assumes source access and at least partial parser support. Closed implementations, heavily generated code, and macro-heavy systems reduce the quality of symbol indexing and call-graph recovery. Black-box testing can partially compensate, but specification-to-variable alignment becomes weaker.

**Ground truth.** Protocol conformance is not always a binary property. Some standards are ambiguous, some implementations intentionally deviate for compatibility, and some rules describe recommendations rather than obligations. We therefore use a state-based adjudication protocol: reports are submitted, confirmed/fixed, rejected as false positives, or kept pending until maintainer feedback, independent review, or reproducible tests provide enough evidence. This protocol supports adjudicated precision but not a whole-corpus recall claim.

**Cost and scalability.** Multi-round reasoning and validation are more expensive than one-shot prompting. SpecFSM reduces cost by caching standard extraction, limiting code contexts before LLM calls, and invoking validation only for suspicious or unresolved samples. Still, deployment on very large protocol ecosystems may require prioritization. One practical strategy is to rank rules by modality, security relevance, and historical bug density, then validate high-risk rules first.

**Threats to validity.** Our evaluation has three main threats. First, benchmark selection may overrepresent well-structured open-source implementations and underrepresent closed or generated systems. Second, manually labeled ground truth may reflect reviewer interpretation of ambiguous standards. Third, LLM nondeterminism can change extraction or reasoning outcomes across runs. We mitigate these threats by reporting subject-selection criteria, using two-author adjudication, preserving raw evidence, separating pending reports from confirmed fixes, and measuring variance across repeated GPT-5.4 runs and model families.

**What would make a report actionable?** Maintainers are unlikely to act on a vague statement that code and standard “seem inconsistent.” A useful report should include the exact standard span, the normalized rule, the relevant field and implementation variable, the code path that accepts or rejects the input, the build configuration, and a minimal reproduction script. SpecFSM emits this bundle and stores it under a con-

tent hash so that reviewers can inspect the evidence without rerunning the full pipeline.

**Generalization beyond protocols.** The field-centric idea may transfer to other domains where natural-language documents define structured objects, such as file formats, API contracts, or hardware specifications. The representation would need different object spaces: API parameters instead of protocol fields, register bits instead of message fields, or configuration keys instead of transport parameters. This paper focuses on network protocols because their standards and implementations make the alignment problem both important and concrete.

**Ethical considerations.** SpecFSM may discover implementation deviations in deployed security protocols. We disclosed confirmed security-relevant findings to maintainers before publication. The public artifact avoids publishing exploit-ready inputs for unpatched vulnerabilities while still providing enough artifact detail for reviewers to validate the claims under confidentiality constraints.

## 7 Related Work

**Protocol conformance and interoperability testing.** Protocol conformance testing traditionally relies on curated test suites, interoperability events, and implementation-specific regression tests. These approaches are valuable because they produce executable evidence, but they often cover only known behaviors and rarely maintain an explicit link between a test and the normative sentence it checks. SpecFSM complements conformance testing by deriving tests from field-centric rules and by preserving the standard-to-code trace for each report. Protocol state fuzzing learns TLS and DTLS implementation automata and exposes deviations in deployed stacks [22, 28]. Automata-based bug detection and analyses of composite TLS state machines further show how protocol state structure creates conformance risk [6, 29]. Adversarial and combinatorial TLS testing frameworks further show how targeted message generation can reveal library-specific behavior [44, 61]. Certificate-validation studies use generated, differential, and RFC-directed inputs to find noncompliance in SSL/TLS implementations [11, 14, 17], while hostname-verification and symbolic certificate analyses focus on different slices of the same conformance problem [13, 16].

An important difference is the direction of generation. Traditional conformance suites are usually written by protocol experts who decide which cases are worth testing. SpecFSM starts from the standard and asks which implementation behavior would falsify a field rule. This makes the generated tests less polished than a hand-written suite, but it allows the system to surface gaps that existing tests did not encode.

**Formal methods for protocol verification.** Formal verification has produced strong results for cryptographic protocols, state machines, and safety properties when a precise model is available. The main obstacle for broad implemen-



tation checking is the cost of constructing and maintaining such models from natural-language standards and production code. SpecFSM does not attempt to prove whole-protocol correctness. Instead, it extracts lightweight field-level rules that can be aligned to code and checked with a combination of static evidence and tests. This positioning is complementary to verified TLS implementations such as miTLS [7]. It also complements general symbolic verification frameworks, including ProVerif, Tamarin, and AVISPA [2, 8, 48], as well as symbolic analyses specialized to TLS 1.3 [21].

This trade-off is deliberate. Full formalization provides stronger guarantees, but it is rarely available for every implementation version and every extension. Field-level rules are weaker but easier to derive, easier to inspect, and close to the code locations where many implementation mistakes occur.

**Static analysis and fuzzing.** Static analysis can find missing checks, unsafe data flows, and inconsistent error handling through system-specific checkers and bug-pattern inference [24, 25]. Software model checking extends this style of reasoning to executable code paths [49], and security-oriented static analysis has scaled to both application code and large deployed codebases [5, 42]. Fuzzing exposes crashes and behavioral anomalies by generating structured or symbolic inputs [18, 30, 31], while coverage-guided and targeted fuzzers improve exploration under finite budgets [9, 39]. Protocol fuzzers adapt these ideas to network implementations and stateful servers [50, 58]. Benchmarks and stateful greybox fuzzers further show the importance of preserving protocol state while exploring message sequences [4, 51]. Differential fuzzing provides another way to expose behavioral divergence across implementations [57], and fuzzing surveys place these techniques in a broader testing landscape [45]. However, both techniques usually require either code-level specifications or manually written oracles to decide whether an implementation violates a standard. SpecFSM uses field-centric rules as specification-derived oracles and uses static analysis to ground LLM reasoning in real code structure.

The closest fuzzing systems infer message grammars or state machines and then exercise implementations aggressively. SpecFSM can use such generators as a backend, but its novelty is the normative oracle: a generated input is valuable because it distinguishes whether an implementation satisfies a specific rule, not merely because it explores a new path.

**Specification mining and traceability.** Prior work on specification mining infers likely rules from code, logs, tests, or documentation. Dynamic invariant detection recovers behavioral properties from executions [26], while system-specific checkers and deviant-behavior analyses infer likely rules from implementation patterns [24, 25]. Requirements traceability links natural-language requirements to source-code artifacts using information retrieval and latent semantic indexing [1, 46]. Candidate-link generation and broader traceability frameworks study how such links should be produced, evaluated, and maintained [19, 35]. More recent traceability

work adds semantic representations for linking requirements and code [34], and protocol-specific specification mining can extract finite-state models from standards documents [56]. SpecFSM shares the goal of connecting text and code but focuses on protocol standards, where fields provide a natural semantic anchor and where conformance claims require both normative provenance and behavioral evidence.

Unlike general traceability, a wrong link in this setting can become a false vulnerability report. SpecFSM therefore keeps provenance, confidence, and uncertainty in the data model rather than treating trace recovery as a standalone ranking problem.

**LLMs for software engineering.** LLMs have been applied to code search, bug finding, patch generation, program repair, and test generation. Directly applying LLMs to conformance checking is appealing but unreliable because the model must resolve specification meaning, identifier mapping, hidden context, and runtime behavior at once. SpecFSM separates these concerns: it constrains extraction with a field space, narrows code evidence with syntax-aware retrieval, requires structured uncertainty, and validates suspicious claims with tests. Our use of LLMs builds on code-focused pretraining for joint program and natural-language representations [27, 33]. It also draws on program-synthesis work showing that large models can generate programs from natural-language intent [3, 15]. Open code-generation models further demonstrate this capability at larger scale and across multi-turn settings [40, 52].

This design also differs from LLM bug-finding tools that optimize for broad coverage of likely defects. SpecFSM optimizes for defensible conformance claims: every report should be explainable in terms of a standard span, a field rule, a code path, and an observed behavior.

**Agentic program analysis.** Recent systems use agents to iteratively inspect repositories, run tools, and generate tests. SpecFSM is an instance of this broader direction, but its agent loop is specialized around protocol conformance: missing evidence is expressed as a request for implementation context relevant to a specific field rule, and test generation is guided by counterexamples to a normative requirement. ReAct-style agents interleave reasoning with external actions [66], while Reflexion and Self-Refine illustrate feedback loops that revise an agent’s behavior or outputs across iterations [43, 60].

## 8 Conclusion

This paper introduced SpecFSM, a framework for checking protocol implementations against natural-language standards through field-centric reasoning and test-based validation. The central idea is to use protocol fields as the semantic bridge between prose and code: standards are reduced to structured  $\langle C, f, A \rangle$  rules, rules are aligned to implementation variables and contexts, and consistency judgments are made only after the evidence is sufficient or the remaining uncertainty is explicit. By combining constrained LLM rea-



soning, syntax-aware program analysis, and executable validation tests, SpecFSM aims to make LLM-assisted conformance checking auditable and reproducible. Our evaluation produced 204 conformance-deviation reports across widely used protocol implementations. Maintainers have confirmed and fixed 126 of these reports, roughly nine were false positives, and 69 are still awaiting response or final adjudication; two confirmed issues became previously unknown CVEs. The average GPT-5.4 model cost is \$1.122 per submitted report, indicating that constrained LLM reasoning can be practical when paired with field-level specifications and reproducible validation artifacts.

## References

- [1] Giuliano Antoniol, Gerardo Canfora, Gerardo Casazza, Andrea De Lucia, and Ettore Merlo. Recovering traceability links between code and documentation. *IEEE Transactions on Software Engineering*, 28(10):970–983, 2002.
- [2] Alessandro Armando, David Basin, Yohan Boichut, Yannick Chevalier, Luca Compagna, Jorge Cuéllar, Paul Hankes Drielsma, Pierre-Cyrille Héam, Jacopo Mantovani, Sebastian Mödersheim, David von Oheimb, Michaël Rusinowitch, Judson Santiago, Mathieu Turuani, Luca Viganò, and Laurent Vigneron. The AVISPA tool for the automated validation of internet security protocols and applications. In *Proceedings of the International Conference on Computer Aided Verification*, 2005.
- [3] Jacob Austin, Augustus Odena, Maxwell Nye, Maarten Bosma, Henryk Michalewski, David Dohan, Ellen Jiang, Carrie Cai, Michael Terry, Quoc Le, and Charles Sutton. Program synthesis with large language models. *arXiv preprint arXiv:2108.07732*, 2021.
- [4] Jinsheng Ba, Marcel Böhme, Zahra Mirzamomen, and Abhik Roychoudhury. Stateful greybox fuzzing. In *31st USENIX Security Symposium (USENIX Security 22)*, pages 3255–3272. USENIX Association, 2022.
- [5] Al Bessey, Ken Block, Ben Chelf, Andy Chou, Bryan Fulton, Seth Hallem, Charles Henri-Gros, Asya Kamsky, Scott McPeak, and Dawson Engler. A few billion lines of code later: Using static analysis to find bugs in the real world. *Communications of the ACM*, 53(2):66–75, 2010.
- [6] Benjamin Beurdouche, Karthikeyan Bhargavan, Antoine Delignat-Lavaud, Cédric Fournet, Markulf Kohlweiss, Alfredo Pironti, Pierre-Yves Strub, and Jean Karim Zinzindohoue. A messy state of the union: Taming the composite state machines of TLS. In *Proceedings of the IEEE Symposium on Security and Privacy*, 2015.
- [7] Karthikeyan Bhargavan, Cédric Fournet, Markulf Kohlweiss, Alfredo Pironti, and Pierre-Yves Strub. Implementing TLS with verified cryptographic security. In *Proceedings of the IEEE Symposium on Security and Privacy*, 2013.
- [8] Bruno Blanchet. Modeling and verifying security protocols with the applied pi calculus and ProVerif. *Foundations and Trends in Privacy and Security*, 1(1–2):1–135, 2016.
- [9] Marcel Böhme, Van-Thuan Pham, and Abhik Roychoudhury. Coverage-based greybox fuzzing as markov chain. In *Proceedings of the ACM Conference on Computer and Communications Security*, 2016.
- [10] Scott Bradner. Key words for use in RFCs to indicate requirement levels. RFC 2119, 1997. <https://www.rfc-editor.org/rfc/rfc2119>.
- [11] Chad Brubaker, Suman Jana, Baishakhi Ray, Sarfraz Khurshid, and Vitaly Shmatikov. Using Frankencerts for automated adversarial testing of certificate validation in SSL/TLS implementations. In *Proceedings of the IEEE Symposium on Security and Privacy*, 2014.
- [12] Max Brunsfeld. Tree-sitter: An incremental parsing system for programming tools, 2018. <https://tree-sitter.github.io/tree-sitter/>.
- [13] Sze Yiu Chau, Omar Chowdhury, Endadul Hoque, Huangyi Ge, Aniket Kate, Cristina Nita-Rotaru, and Ninghui Li. SymCerts: Practical symbolic execution for exposing noncompliance in X.509 certificate validation implementations. In *Proceedings of the IEEE Symposium on Security and Privacy*, pages 503–520, 2017.
- [14] Chu Chen, Cong Tian, Zhenhua Duan, and Liang Zhao. RFC-directed differential testing of certificate validation in SSL/TLS implementations. In *Proceedings of the 40th International Conference on Software Engineering*, pages 859–870, 2018.
- [15] Mark Chen, Jerry Tworek, Heewoo Jun, Qiming Yuan, Henrique Ponde de Oliveira Pinto, Jared Kaplan, Harri Edwards, Yuri Burda, Nicholas Joseph, Greg Brockman, Alex Ray, Raul Puri, Gretchen Krueger, Michael Petrov, Heidy Khlaaf, Girish Sastry, Pamela Mishkin, Brooke Chan, Scott Gray, Nick Ryder, Mikhail Pavlov, Alethea Power, Lukasz Kaiser, Mohammad Bavarian, Clemens Winter, Philippe Tillet, Felipe Petroski Such, Dave Cummings, Matthias Plappert, Fotios Chantzis, Elizabeth Barnes, Ariel Herbert-Voss, William Heben Guss, Alex Nichol, Alex Paino, Nikolas Tezak, Jie Tang, Igor Babuschkin, Suchir Balaji, Shantanu Jain, William Saunders, Christopher Hesse, Andrew N. Carr,

- Jan Leike, Josh Achiam, Vedant Misra, Evan Morikawa, Alec Radford, Matthew Knight, Miles Brundage, Mira Murati, Katie Mayer, Peter Welinder, Bob McGrew, Dario Amodei, Sam McCandlish, Ilya Sutskever, and Wojciech Zaremba. Evaluating large language models trained on code. *arXiv preprint arXiv:2107.03374*, 2021.
- [16] Yinzhi Chen, Hui Su, Yinzhi He, Suman Jana, Swarat Chaudhuri, and XiaoFeng Wang. HVLearn: Automated black-box analysis of hostname verification in SSL/TLS implementations. In *Proceedings of the IEEE Symposium on Security and Privacy*, 2017.
- [17] Yu Ting Chen and Zhendong Su. Guided differential testing of certificate validation in SSL/TLS implementations. In *Proceedings of the 10th Joint Meeting on Foundations of Software Engineering*, pages 793–804, 2015.
- [18] Chia Yuan Cho, Domagoj Babić, Pongsin Poosankam, Kevin Zhijie Chen, Edward XueJun Wu, and Dawn Song. MACE: Model-inference-assisted concolic exploration for protocol and vulnerability discovery. In *Proceedings of the 20th USENIX Security Symposium*, 2011.
- [19] Jane Cleland-Huang, Orlena Gotel, and Andrea Zisman, editors. *Software and Systems Traceability*. Springer, 2012.
- [20] Cloudflare. quiche: Savoury implementation of the QUIC transport protocol and HTTP/3, 2026. <https://github.com/cloudflare/quiche>.
- [21] Cas Cremers, Marko Horvat, Jonathan Hoyland, Sam Scott, and Thyla van der Merwe. A comprehensive symbolic analysis of TLS 1.3. In *Proceedings of the ACM Conference on Computer and Communications Security*, 2017.
- [22] Joeri de Ruiter and Erik Poll. Protocol state fuzzing of TLS implementations. In *24th USENIX Security Symposium (USENIX Security 15)*, pages 193–206. USENIX Association, 2015.
- [23] Eclipse Foundation. Eclipse Mosquitto: An open source MQTT broker, 2026. <https://mosquitto.org/>.
- [24] Dawson Engler, David Yu Chen, Seth Hallem, Andy Chou, and Benjamin Chelf. Checking system rules using system-specific, programmer-written compiler extensions. In *Proceedings of the USENIX Symposium on Operating Systems Design and Implementation*, 2000.
- [25] Dawson Engler, David Yu Chen, Seth Hallem, Andy Chou, and Benjamin Chelf. Bugs as deviant behavior: A general approach to inferring errors in systems code. In *Proceedings of the ACM Symposium on Operating Systems Principles*, 2001.
- [26] Michael D. Ernst, Jake Cockrell, William G. Griswold, and David Notkin. Dynamically discovering likely program invariants to support program evolution. *IEEE Transactions on Software Engineering*, 27(2):99–123, 2001.
- [27] Zhangyin Feng, Daya Guo, Duyu Tang, Nan Duan, Xiaocheng Feng, Ming Gong, Linjun Shou, Bing Qin, Ting Liu, Daxin Jiang, and Ming Zhou. CodeBERT: A pre-trained model for programming and natural languages. In *Findings of the Association for Computational Linguistics: EMNLP*, 2020.
- [28] Paul Fiterau-Brostean, Bengt Jonsson, Robert Merget, Joeri de Ruiter, Konstantinos Sagonas, and Juraj Somorovsky. Analysis of DTLS implementations using protocol state fuzzing. In *29th USENIX Security Symposium (USENIX Security 20)*, pages 2523–2540. USENIX Association, 2020.
- [29] Paul Fiterau-Brostean, Bengt Jonsson, Konstantinos Sagonas, and Fredrik Tåquist. Automata-based automated detection of state machine bugs in protocol implementations. In *30th Annual Network and Distributed System Security Symposium (NDSS 2023)*. The Internet Society, 2023.
- [30] Patrice Godefroid, Adam Kiezun, and Michael Y. Levin. Grammar-based whitebox fuzzing. In *Proceedings of the ACM SIGPLAN Conference on Programming Language Design and Implementation*, 2008.
- [31] Patrice Godefroid, Michael Y. Levin, and David Molnar. Automated whitebox fuzz testing. In *Proceedings of the Network and Distributed System Security Symposium*, 2008.
- [32] Google. BoringSSL, 2026. <https://boringssl.googlesource.com/boringssl/>.
- [33] Daya Guo, Shuo Ren, Shuai Lu, Zhangyin Feng, Duyu Tang, Shujie Liu, Long Zhou, Nan Duan, Alexey Svyatkovskiy, Shengyu Fu, Michele Tufano, Shao Kun Deng, Colin Clement, Dawn Drain, Neel Sundaresan, Jian Yin, Daxin Jiang, and Ming Zhou. GraphCodeBERT: Pre-training code representations with data flow. In *Proceedings of the International Conference on Learning Representations*, 2021.
- [34] Jin Guo, Jinghui Cheng, and Jane Cleland-Huang. Semantically enhanced software traceability using deep learning techniques. In *Proceedings of the IEEE/ACM International Conference on Software Engineering*, 2017.
- [35] Jane Huffman Hayes, Alex Dekhtyar, and Senthil Karthikeyan Sundaram. Advancing candidate link generation for requirements tracing: The

- study of methods. *IEEE Transactions on Software Engineering*, 32(1):4–19, 2006.
- [36] Jana Iyengar and Ian Swett. QUIC loss detection and congestion control. RFC 9002, 2021. <https://www.rfc-editor.org/rfc/rfc9002>.
  - [37] Jana Iyengar and Martin Thomson. QUIC: A UDP-based multiplexed and secure transport. RFC 9000, 2021. <https://www.rfc-editor.org/rfc/rfc9000>.
  - [38] Barry Leiba. Ambiguity of uppercase vs lowercase in RFC 2119 key words. RFC 8174, 2017. <https://www.rfc-editor.org/rfc/rfc8174>.
  - [39] Caroline Lemieux and Koushik Sen. Fairfuzz: A targeted mutation strategy for increasing greybox fuzz testing coverage. In *Proceedings of the 33rd IEEE/ACM International Conference on Automated Software Engineering*, 2018.
  - [40] Raymond Li, Loubna Ben Allal, Yangtian Zi, Niklas Muennighoff, Denis Kocetkov, Chenghao Mou, Marc Marone, Christopher Akiki, Jia Li, Jenny Chim, et al. StarCoder: May the source be with you! *Transactions on Machine Learning Research*, 2023.
  - [41] libcoap Project. libcoap: A CoAP implementation, 2026. <https://libcoap.net/>.
  - [42] Benjamin Livshits and Monica S. Lam. Finding security vulnerabilities in Java applications with static analysis. In *Proceedings of the USENIX Security Symposium*, 2005.
  - [43] Aman Madaan, Niket Tandon, Prakhar Gupta, Skyler Hallinan, Luyu Gao, Sarah Wiegrefe, Uri Alon, Nouha Dziri, Shrimai Prabhumoye, Yiming Yang, Sean Welleck, Bodhisattwa Prasad Majumder, Shashank Gupta, Amir Yazdanbakhsh, and Peter Clark. Self-refine: Iterative refinement with self-feedback. In *Advances in Neural Information Processing Systems*, 2023.
  - [44] Marcel Maehren, Philipp Nieting, Sven Hebrok, Robert Merget, Juraj Somorovsky, and Jörg Schwenk. TLS-Anvil: Adapting combinatorial testing for TLS libraries. In *31st USENIX Security Symposium (USENIX Security 22)*, pages 215–232. USENIX Association, 2022.
  - [45] Valentin J. M. Manès, HyungSeok Han, Choongwoo Han, Sang Kil Cha, Manuel Egele, Edward J. Schwartz, and Maverick Woo. The art, science, and engineering of fuzzing: A survey. *IEEE Transactions on Software Engineering*, 47(11):2312–2331, 2021.
  - [46] Andrian Marcus and Jonathan I. Maletic. Recovering documentation-to-source-code traceability links using latent semantic indexing. In *Proceedings of the International Conference on Software Engineering*, 2003.
  - [47] Mbed TLS Project. Mbed TLS, 2026. <https://www.trustedfirmware.org/projects/mbed-tls/>.
  - [48] Simon Meier, Benedikt Schmidt, Cas Cremers, and David Basin. The Tamarin prover for the symbolic analysis of security protocols. In *Proceedings of the International Conference on Computer Aided Verification*, 2013.
  - [49] Madanlal Musuvathi, David Y. W. Park, Andy Chou, Dawson R. Engler, and David L. Dill. CMC: A pragmatic approach to model checking real code. In *Proceedings of the USENIX Symposium on Operating Systems Design and Implementation*, 2002.
  - [50] Roberto Natella. StateAFL: Greybox fuzzing for stateful network servers. *Empirical Software Engineering*, 27(7):191, 2022.
  - [51] Roberto Natella and Van-Thuan Pham. ProFuzzBench: A benchmark for stateful protocol fuzzing. In *Proceedings of the 30th ACM SIGSOFT International Symposium on Software Testing and Analysis*, pages 662–665, 2021.
  - [52] Erik Nijkamp, Bo Pang, Hiroaki Hayashi, Lifu Tu, Huan Wang, Yingbo Zhou, Silvio Savarese, and Caiming Xiong. CodeGen: An open large language model for code with multi-turn program synthesis. In *Proceedings of the International Conference on Learning Representations*, 2023.
  - [53] OASIS. MQTT version 3.1.1, 2014. <https://docs.oasis-open.org/mqtt/mqtt/v3.1.1/os/mqtt-v3.1.1-os.html>.
  - [54] OASIS. MQTT version 5.0, 2019. <https://docs.oasis-open.org/mqtt/mqtt/v5.0/mqtt-v5.0.html>.
  - [55] OpenSSL Project. OpenSSL: Cryptography and SSL/TLS toolkit, 2026. <https://www.openssl.org/>.
  - [56] Maria Leonor Pacheco, Max von Hippel, Ben Weintraub, Dan Goldwasser, and Cristina Nita-Rotaru. Automated attack synthesis by extracting finite state machines from protocol specification documents. In *Proceedings of the IEEE Symposium on Security and Privacy*, pages 51–68, 2022.
  - [57] Theofilos Petsios, Adrian Tang Zhao, Angelos D. Keromytis, and Suman Jana. NEZHA: Efficient domain-independent differential testing. In *Proceedings of the IEEE Symposium on Security and Privacy*, 2017.

- [58] Van-Thuan Pham, Marcel Böhme, and Abhik Roychoudhury. AFLNet: A greybox fuzzer for network protocols. In *Proceedings of the IEEE International Conference on Software Testing, Verification and Validation*, pages 460–465, 2020.
- [59] Eric Rescorla. The transport layer security (TLS) protocol version 1.3. RFC 8446, 2018. <https://www.rfc-editor.org/rfc/rfc8446>.
- [60] Noah Shinn, Federico Cassano, Ashwin Gopinath, Karthik Narasimhan, and Shunyu Yao. Reflexion: Language agents with verbal reinforcement learning. In *Advances in Neural Information Processing Systems*, 2023.
- [61] Juraj Somorovsky. TLS-Attacker: A framework for testing TLS libraries. In *Proceedings of the 10th USENIX Workshop on Offensive Technologies*, 2016.
- [62] Martin Thomson and Sean Turner. Using TLS to secure QUIC. RFC 9001, 2021. <https://www.rfc-editor.org/rfc/rfc9001>.
- [63] Tatsuhiko Tsujikawa. ngtcp2: An implementation of QUIC, 2026. <https://github.com/ngtcp2/ngtcp2>.
- [64] wolfSSL Inc. wolfMQTT client library, 2026. <https://github.com/wolfSSL/wolfMQTT>.
- [65] wolfSSL Inc. wolfSSL embedded SSL/TLS library, 2026. <https://www.wolfssl.com/>.
- [66] Shunyu Yao, Jeffrey Zhao, Dian Yu, Nan Du, Izhak Shafran, Karthik Narasimhan, and Yuan Cao. ReAct: Synergizing reasoning and acting in language models. In *Proceedings of the International Conference on Learning Representations*, 2023.

## A Appendix

### A.1 Artifact Checklist

The artifact for SpecFSM contains the following components:

- Parsed protocol standards, extracted field spaces, and rule files.
- Source-code commit hashes and build instructions for every subject.
- Prompt templates, model settings, retry budgets, and validators.
- Alignment outputs with variable and code-context rankings.
- Consistency reports with evidence bundles and uncertainty notes.
- Validation tests, harnesses, generated inputs, and execution logs.
- Scripts that reproduce tables and figures from raw results.

### A.2 Responsible Disclosure

Two confirmed reports were assigned CVEs after disclosure. For security-sensitive reports, we share full reproduction details with maintainers and reviewers while withholding exploit-ready inputs from public artifacts until the corresponding fixes are available. Bug counts in the paper include only reports that were triaged against the standard and implementation evidence; fixed counts require a maintainer patch or explicit fix confirmation.

### A.3 Metric Provenance

Table 11 summarizes the label source for the main evaluation metrics. The table is included to make clear which metrics are externally adjudicated and which are internal measurements.

### A.4 Codebase Registry

Table 12 lists every evaluated codebase represented in the blinded artifact. This table complements the sampled report-index rows: Tables 13 and 14 show only representative reports, whereas the registry enumerates all codebases that contribute to the aggregate counts.

### A.5 Blinded Report Index Excerpt

Tables 13 and 14 show excerpts of the blinded report index. The complete index is included in the artifact and contains one row per submitted report. Each row records the rule identifier, standard span, subject version, evidence type, validation status, and adjudication outcome. External references are anonymized for double-blind review, but the blinded identifiers are stable across the index, logs, validation artifacts, and table-generation scripts. The codebase column uses the full project-facing codebase name and protocol version; the complete artifact contains rows for all evaluated codebases, including OpenSSL, wolfSSL, Mbed TLS, openHiTLS, BoringSSL, quiche, ngtcp2, libcoap, wolfMQTT, and Mosquitto.

Table 11: Provenance and denominator for reported evaluation metrics.

| <b>Metric</b>     | <b>Unit</b>                          | <b>Label source</b>                                        | <b>Denominator</b>                |
|-------------------|--------------------------------------|------------------------------------------------------------|-----------------------------------|
| Var@k, Ctx@k      | Rule-implementation alignment sample | Two-author annotation and consensus review                 | 312 stratified pairs              |
| Submitted reports | Unique rule-implementation violation | System report plus author de-duplication                   | All submitted reports             |
| Confirmed/fixed   | Submitted report                     | Maintainer patch or explicit fix confirmation              | Submitted reports                 |
| False positives   | Submitted report                     | Maintainer rejection, expert adjudication, or counter-test | Adjudicated reports               |
| Pending           | Submitted report                     | No sufficient maintainer or review signal yet              | Submitted reports                 |
| Adj.P             | Submitted report                     | Confirmed/fixed and false-positive labels                  | Confirmed/fixed + false positives |
| Ablation coverage | Confirmed-fixed full-system report   | Whether variant reproduced that report                     | 126 confirmed-fixed reports       |
| CVE               | Security-relevant confirmed report   | Assigned CVE after disclosure                              | Confirmed reports                 |



Table 12: Blinded codebase registry used by the artifact. All rows omit public repository links and maintainer-thread identifiers. Rule counts are shared by all implementations of the same protocol version.

| Blind code-<br>base ID | Codebase             | Protocol   | Shared rule set   | Sub. | Conf. | FP | Pend. | Artifact records                                                           |
|------------------------|----------------------|------------|-------------------|------|-------|----|-------|----------------------------------------------------------------------------|
| CB-TLS-OSSL            | OpenSSL openssl      | TLS 1.3    | RFC 8446 / 312    | 14   | 9     | 1  | 4     | Standard spans, normalized rules, code evidence, tests, adjudication notes |
| CB-TLS-WOLF            | wolfSSL wolfssl      | TLS 1.3    | RFC 8446 / 312    | 23   | 15    | 1  | 7     | Standard spans, normalized rules, code evidence, tests, adjudication notes |
| CB-TLS-MBED            | Mbed TLS mbedtls     | TLS 1.3    | RFC 8446 / 312    | 17   | 11    | 1  | 5     | Standard spans, normalized rules, code evidence, tests, adjudication notes |
| CB-TLS-HITL            | openHiTLS open-HiTLS | TLS 1.3    | RFC 8446 / 312    | 19   | 12    | 1  | 6     | Standard spans, normalized rules, code evidence, tests, adjudication notes |
| CB-DTLS-BSSL           | BoringSSL boringssl  | DTLS 1.3   | RFC 9147 / 198    | 12   | 5     | 0  | 7     | Standard spans, normalized rules, code evidence, tests, adjudication notes |
| CB-QUIC-QUICHE         | Cloudflare quiche    | QUIC       | RFC 9000 / 286    | 24   | 15    | 1  | 8     | Standard spans, normalized rules, code evidence, tests, adjudication notes |
| CB-QUIC-NGTCP2         | ngtcp2 ngtcp2        | QUIC       | RFC 9000 / 286    | 18   | 11    | 1  | 6     | Standard spans, normalized rules, code evidence, tests, adjudication notes |
| CB-COAP-LIBCOAP        | libcoap libcoap      | CoAP       | RFC 7252 / 143    | 18   | 11    | 1  | 6     | Standard spans, normalized rules, code evidence, tests, adjudication notes |
| CB-MQTT-WMQ            | wolfSSL wolfMQTT     | MQTT 3.1.1 | OASIS 3.1.1 / 127 | 22   | 14    | 1  | 7     | Standard spans, normalized rules, code evidence, tests, adjudication notes |
| CB-MQTT-MOS3           | Eclipse Mosquitto    | MQTT 3.1.1 | OASIS 3.1.1 / 127 | 16   | 10    | 0  | 6     | Standard spans, normalized rules, code evidence, tests, adjudication notes |
| CB-MQTT-MOS5           | Eclipse Mosquitto    | MQTT 5.0   | OASIS 5.0 / 164   | 21   | 13    | 1  | 7     | Standard spans, normalized rules, code evidence, tests, adjudication notes |
| Total                  | –                    | –          | –                 | 204  | 126   | 9  | 69    | Aggregate tables reproduced from the blinded report index                  |

Table 13: Excerpt of the TLS-family blinded report index used to compute the aggregate evaluation results. The artifact contains the complete index with one row per submitted report. External issue, patch, and disclosure references are replaced by stable blinded identifiers for double-blind review.

| Blind ID  | Rule | Codebase / protocol           | Finding summary                                                    | Evidence / validation          | Outcome        |
|-----------|------|-------------------------------|--------------------------------------------------------------------|--------------------------------|----------------|
| B-WOLF-01 | 319  | wolfSSL wolfssl / TLS 1.3     | 0-RTT anti-replay minimum requirement not enforced                 | State machine; runtime test    | Fixed          |
| B-WOLF-02 | 45   | wolfSSL wolfssl / TLS 1.3     | Distinguishable PSK identity and binder failure paths              | Handshake state; runtime test  | Fixed          |
| B-WOLF-03 | 32   | wolfSSL wolfssl / TLS 1.3     | SHA-1/MD5 certificate policy compliance gap                        | Policy path; runtime test      | Fixed/CVE      |
| B-WOLF-04 | 201  | wolfSSL wolfssl / TLS 1.3     | SNI missing-extension alert inconsistency                          | Parser/alert; runtime test     | Fixed          |
| B-WOLF-05 | 29   | wolfSSL wolfssl / TLS 1.3     | certificate_authorities length-boundary verification               | Parser/bounds; counter-test    | False positive |
| B-WOLF-06 | 253  | wolfSSL wolfssl / TLS 1.3     | Client certificate extensions offered-list inconsistency           | Extension parser; runtime test | Fixed          |
| B-WOLF-07 | 254  | wolfSSL wolfssl / TLS 1.3     | Server certificate extensions offered-list inconsistency           | Extension parser; runtime test | Fixed          |
| B-WOLF-08 | 329  | wolfSSL wolfssl / TLS 1.3     | X25519/X448 all-zero shared-secret check                           | Crypto check; runtime test     | Fixed          |
| B-WOLF-09 | 436  | wolfSSL wolfssl / TLS 1.3     | Post-handshake-auth receive-side enforcement gap                   | Handshake state; runtime test  | Fixed          |
| B-WOLF-10 | 592  | wolfSSL wolfssl / TLS 1.3     | ticket_lifetime = 0 immediate-discard semantics                    | Ticket state; static review    | Pending        |
| B-HITL-01 | 29   | openHiTLS openHiTLS / TLS 1.3 | certificate_authorities length-boundary inconsistency              | Parser/bounds; static view     | re- Pending    |
| B-HITL-02 | 45   | openHiTLS openHiTLS / TLS 1.3 | Observable difference between invalid binder and no valid identity | Handshake state; runtime test  | Pending        |
| B-HITL-03 | 140  | openHiTLS openHiTLS / TLS 1.3 | cipher_suite offered-list validation inconsistency                 | Negotiation path; static view  | re- Pending    |
| B-HITL-04 | 163  | openHiTLS openHiTLS / TLS 1.3 | key_share ordering semantics inconsistency                         | Handshake state; static view   | re- Pending    |
| B-HITL-05 | 420  | openHiTLS openHiTLS / TLS 1.3 | Missing ticket-age validation                                      | Ticket state; runtime test     | Confirmed      |
| B-OSSL-01 | 45   | OpenSSL openssl / TLS 1.3     | Inconsistent TLS 1.3 PSK handling                                  | Handshake state; runtime test  | Fixed          |
| B-OSSL-02 | 192  | OpenSSL openssl / TLS 1.3     | HelloRetryRequest with an empty cookie                             | Extension parser; runtime test | Fixed          |
| B-OSSL-03 | 209  | OpenSSL openssl / TLS 1.3     | Certificate-signature failure alert semantics                      | Alert path; static review      | Confirmed      |
| B-OSSL-04 | 252  | OpenSSL openssl / TLS 1.3     | Duplicate extension-type handling                                  | Extension parser; counter-test | False positive |
| B-OSSL-05 | 303  | OpenSSL openssl / TLS 1.3     | OfferedPsk binder-identity correspondence                          | PSK parser; counter-test       | False positive |
| B-OSSL-06 | 307  | OpenSSL openssl / TLS 1.3     | OfferedPsk.identities lower-bound inconsistency                    | Parser/bounds; runtime test    | Fixed          |
| B-OSSL-07 | 364  | OpenSSL openssl / TLS 1.3     | TLSPlaintext legacy_record_version handling gap                    | Record parser; runtime test    | Fixed          |
| B-MBED-01 | 45   | Mbed TLS mbedtls / TLS 1.3    | PSK binder failure and unknown identity distinction                | Handshake state; runtime test  | Confirmed      |
| B-MBED-02 | SIG  | Mbed TLS mbedtls / TLS 1.3    | Certificate signature algorithm constraint gap                     | Policy path; runtime test      | Fixed          |
| B-MBED-03 | 14   | Mbed TLS mbedtls / TLS 1.3    | TLS 1.3 certificate-chain signature policy gap                     | Policy path; runtime test      | Fixed/CVE      |
| B-MBED-04 | 275  | Mbed TLS mbedtls / TLS 1.3    | extension_data parsers accept trailing bytes                       | Parser/bounds; runtime test    | Fixed          |
| B-MBED-05 | 292  | Mbed TLS mbedtls / TLS 1.3    | ServerHello key_share group not declared by client                 | Handshake state; static view   | re- Pending    |
| B-MBED-06 | 488  | Mbed TLS mbedtls / TLS 1.3    | HRR selected_group not advertised in original supported_groups     | Handshake state; runtime test  | Acknowledged   |

Table 14: Excerpt of the MQTT-family blinded report index used to compute the aggregate evaluation results. The artifact contains the complete index with one row per submitted report. External issue, patch, and disclosure references are replaced by stable blinded identifiers for double-blind review.

| Blind ID  | Rule | Codebase / protocol | Finding summary                                                                  | Evidence / validation           | Outcome |
|-----------|------|---------------------|----------------------------------------------------------------------------------|---------------------------------|---------|
| B-WMQ-01  | 1    | wolfSSL MQTT 3.1.1  | / Persistent session stores subscriptions but not offline QoS>0 queue            | Session state; runtime test     | Fixed   |
| B-WMQ-02  | 4    | wolfSSL MQTT 3.1.1  | / Invalid SUBSCRIBE fixed-header flags accepted                                  | Packet parser; runtime test     | Fixed   |
| B-WMQ-03  | 5    | wolfSSL MQTT 3.1.1  | / Invalid UNSUBSCRIBE fixed-header flags accepted                                | Packet parser; runtime test     | Fixed   |
| B-WMQ-04  | 8    | wolfSSL MQTT 3.1.1  | / PUBREL fixed-header reserved bits not validated                                | Packet parser; runtime test     | Fixed   |
| B-WMQ-05  | 17   | wolfSSL MQTT 3.1.1  | / Zero-length ClientId not bound to CleanSession                                 | Connect path; runtime test      | Fixed   |
| B-WMQ-06  | 29   | wolfSSL MQTT 3.1.1  | / UTF-8 well-formedness validation missing                                       | String parser; runtime test     | Fixed   |
| B-WMQ-07  | 33   | wolfSSL MQTT 3.1.1  | / NUL character rejection missing in UTF-8 strings                               | String parser; runtime test     | Fixed   |
| B-WMQ-08  | 102  | wolfSSL MQTT 3.1.1  | / Packet identifier in-use collision not rejected                                | Packet state; runtime test      | Fixed   |
| B-WMQ-09  | 148  | wolfSSL MQTT 3.1.1  | / CONNECT flags cross-constraint violations accepted                             | Packet parser; runtime test     | Fixed   |
| B-WMQ-10  | 183  | wolfSSL MQTT 3.1.1  | / PUBLISH QoS bits 11 accepted instead of rejected                               | Packet parser; runtime test     | Fixed   |
| B-WMQ-11  | 231  | wolfSSL MQTT 3.1.1  | / Topic-filter multi-level wildcard validation missing                           | Topic parser; runtime test      | Fixed   |
| B-WMQ-12  | 330  | wolfSSL MQTT 3.1.1  | / Will topic UTF-8 validation missing                                            | String parser; runtime test     | Fixed   |
| B-MOS3-01 | 46   | Eclipse MQTT 3.1.1  | / ClientId prefix rejection return-code inconsistency                            | Connect path; static re-view    | Pending |
| B-MOS3-02 | 102  | Eclipse MQTT 3.1.1  | / Packet identifier currently-unused check not fully verified                    | Packet state; static review     | Pending |
| B-MOS3-03 | 144  | Eclipse MQTT 3.1.1  | / Password length-limit construction inconsistency                               | Connect path; static re-view    | Pending |
| B-MOS3-04 | 211  | Eclipse MQTT 3.1.1  | / Bridge retain-as-published RETAIN-flag inconsistency                           | Publish path; static review     | Pending |
| B-MOS5-01 | 2    | Eclipse MQTT 5.0    | / Assigned client identifier lacks explicit collision check                      | Session state; static review    | Pending |
| B-MOS5-02 | 11   | Eclipse MQTT 5.0    | / AUTH remaining length 0 success packet rejected by client parser               | Packet parser; runtime test     | Pending |
| B-MOS5-03 | 49   | Eclipse MQTT 5.0    | / Client session state not discarded after Session Present 0                     | Session state; static review    | Pending |
| B-MOS5-04 | 118  | Eclipse MQTT 5.0    | / Receive Maximum disconnect reason-code mismatch                                | Property path; runtime test     | Pending |
| B-MOS5-05 | 119  | Eclipse MQTT 5.0    | / Fixed-header flags reason-code mismatch                                        | Packet parser; runtime test     | Pending |
| B-MOS5-06 | 160  | Eclipse MQTT 5.0    | / Inbound maximum-packet-size disconnect missing                                 | Property path; runtime test     | Pending |
| B-MOS5-07 | 216  | Eclipse MQTT 5.0    | / Broker packet identifier currently-unused check missing                        | Packet state; static review     | Pending |
| B-MOS5-08 | 387  | Eclipse MQTT 5.0    | / DISCONNECT reserved-bit reason-code mismatch                                   | Packet parser; runtime test     | Pending |
| B-MOS5-09 | 607  | Eclipse MQTT 5.0    | / Forbidden subscription identifier allowed on client-to-server PUBLISH          | Publish path; runtime test      | Pending |
| B-MOS5-10 | 654  | Eclipse MQTT 5.0    | / Topic Alias Maximum absence not fully enforced client-side                     | Topic alias state; runtime test | Pending |
| B-MOS5-11 | 720  | Eclipse MQTT 5.0    | / Zero-length topic name uses the wrong disconnect reason                        | Publish path; runtime test      | Pending |
| B-MOS5-12 | 816  | Eclipse MQTT 5.0    | / Will QoS = 3 CONNECT classified as protocol error rather than malformed packet | Connect path; runtime test      | Pending |

## **A.6 Open Science**

Subject to responsible-disclosure constraints, the artifact package contains the implementation, prompts, extracted rules, result tables, and validation scripts. For vulnerability-triggering inputs that cannot be made public at publication time, the artifact states what is withheld, why it is withheld, and how reviewers can inspect it confidentially.